

```
In [ ]:
```

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.utils import resample
import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
```

Step 1: Preprocessing

Load and View Data

Drop unnecessary columns pertaining to histogram results, which aren't typically a part of CTG results.

```
In [ ]:
```

```
data = pd.read_csv('/content/drive/MyDrive/DM_Project/fetal_health.csv')
feat = data[['baseline value',
             'accelerations',
             'fetal_movement',
             'uterine_contractions',
             'light_decelerations',
             'severe_decelerations',
             'prolonged_decelerations',
             'abnormal_short_term_variability',
             'mean_value_of_short_term_variability',
             'mean_value_of_long_term_variability',
             'percentage_of_time_with_abnormal_long_term_variability',
             'fetal_health']]
feat = feat.rename(columns={'percentage_of_time_with_abnormal_long_term_variability': 'per
cent_time_abnormal_long_variability'})
data.shape
```

```
Out[ ]:
```

```
(2126, 22)
```

```
In [ ]:
```

```
X = feat[['baseline value',
          'accelerations',
          'fetal_movement',
          'uterine_contractions',
          'light_decelerations',
          'severe_decelerations',
          'prolonged_decelerations',
          'abnormal_short_term_variability',
          'mean_value_of_short_term_variability',
          'mean_value_of_long_term_variability',
          'percent_time_abnormal_long_variability']]
y = feat[['fetal_health']]
```

```
In [ ]:
```

```
feat.describe().T
```

```
Out[ ]:
```

	count	mean	std	min	25%	50%	75%	max
baseline value	2126.0	133.303857	9.840844	106.0	126.000	133.000	140.000	160.000
accelerations	2126.0	0.003178	0.003866	0.0	0.000	0.002	0.006	0.019

	fetal_movement	count	mean	std	min	25%	50%	75%	max
	uterine_contractions	2126.0	0.004366	0.002946	0.0	0.002	0.004	0.007	0.015
	light_decelerations	2126.0	0.001889	0.002960	0.0	0.000	0.000	0.003	0.015
	severe_decelerations	2126.0	0.000003	0.000057	0.0	0.000	0.000	0.000	0.001
	prolongued_decelerations	2126.0	0.000159	0.000590	0.0	0.000	0.000	0.000	0.005
	abnormal_short_term_variability	2126.0	46.990122	17.192814	12.0	32.000	49.000	61.000	87.000
	mean_value_of_short_term_variability	2126.0	1.332785	0.883241	0.2	0.700	1.200	1.700	7.000
	mean_value_of_long_term_variability	2126.0	8.187629	5.628247	0.0	4.600	7.400	10.800	50.700
	percent_time_abnormal_long_variability	2126.0	9.846660	18.396880	0.0	0.000	0.000	11.000	91.000
	fetal_health	2126.0	1.304327	0.614377	1.0	1.000	1.000	1.000	3.000

In []:

```
feat.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2126 entries, 0 to 2125
Data columns (total 12 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   baseline value                        2126 non-null   float64
 1   accelerations                        2126 non-null   float64
 2   fetal_movement                       2126 non-null   float64
 3   uterine_contractions                 2126 non-null   float64
 4   light_decelerations                 2126 non-null   float64
 5   severe_decelerations                2126 non-null   float64
 6   prolonged_decelerations              2126 non-null   float64
 7   abnormal_short_term_variability      2126 non-null   float64
 8   mean_value_of_short_term_variability 2126 non-null   float64
 9   mean_value_of_long_term_variability  2126 non-null   float64
10   percent_time_abnormal_long_variability 2126 non-null   float64
11   fetal_health                         2126 non-null   float64
dtypes: float64(12)
memory usage: 199.4 KB
```

Dimensionality Reduction - Principal Component Analysis

In []:

```
# Scale data with standard scalar
sc = StandardScaler().set_output(transform='pandas')
scaled = sc.fit(X).transform(X)
```

In []:

```
pca = PCA()
pca.fit(scaled)
d = {'Feature':scaled.columns.values, 'PCA Variance Ratio':pca.explained_variance_ratio_}
pca_info = pd.DataFrame(data=d)
l=[]

for z in range(1, len(pca_info['PCA Variance Ratio']) + 1):
    l.append(sum(pca_info['PCA Variance Ratio'].iloc[:z]))

pca_info['Sum PCA Variance'] = l
display(pca_info)
```

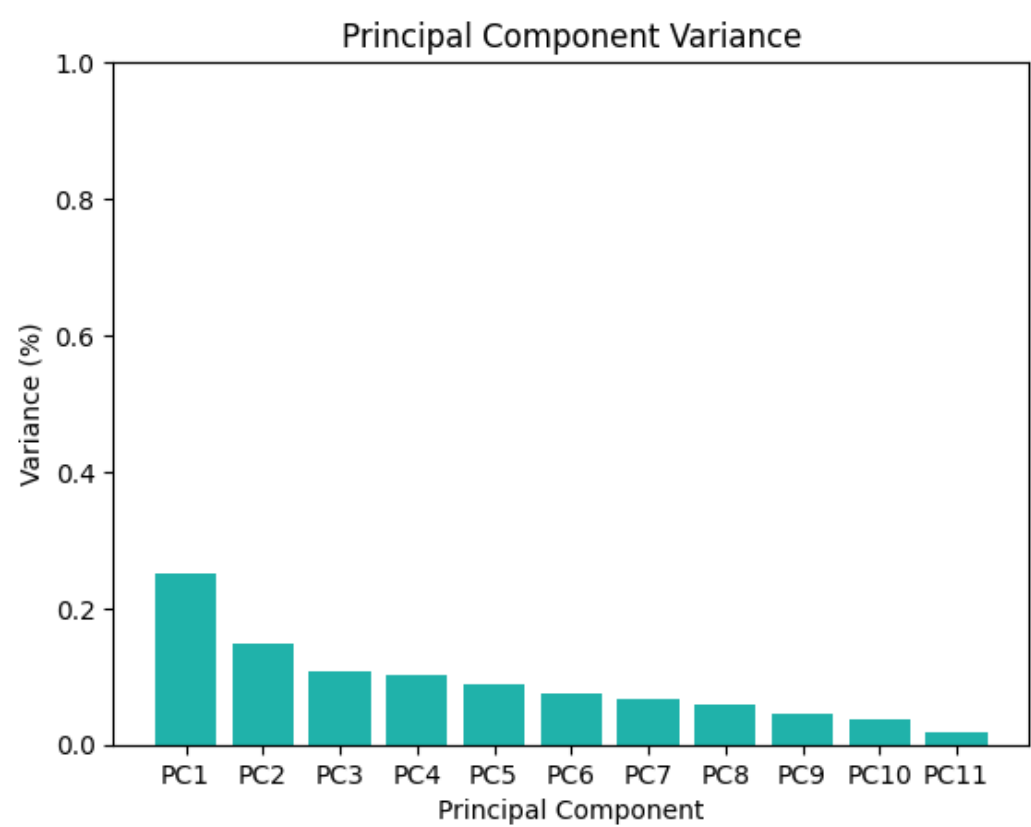
	Feature	PCA Variance Ratio	Sum PCA Variance
0	baseline value	0.250640	0.250640
1	accelerations	0.148219	0.398859

2	fetal_movement	0.107775	0.506633
	Feature	PCA Variance Ratio	Sum PCA Variance
3	uterine_contractions	0.102701	0.609334
4	light_decelerations	0.089323	0.698657
5	severe_decelerations	0.076052	0.774709
6	prolongued_decelerations	0.066112	0.840821
7	abnormal_short_term_variability	0.059541	0.900362
8	mean_value_of_short_term_variability	0.043945	0.944307
9	mean_value_of_long_term_variability	0.036622	0.980929
10	percent_time_abnormal_long_variability	0.019071	1.000000

In []:

```
trained_pca = pca.transform(scaled)
var_ratio = pca_info['PCA Variance Ratio']

x=range(0, len(var_ratio))
plt.bar(x, var_ratio, color='lightseagreen')
plt.ylabel('Variance (%)')
plt.xlabel('Principal Component')
plt.xticks(x, ['PC1', 'PC2', 'PC3', 'PC4', 'PC5', 'PC6', 'PC7', 'PC8', 'PC9', 'PC10', 'P
C11'])
plt.ylim([0,1])
plt.title('Principal Component Variance')
plt.show()
```



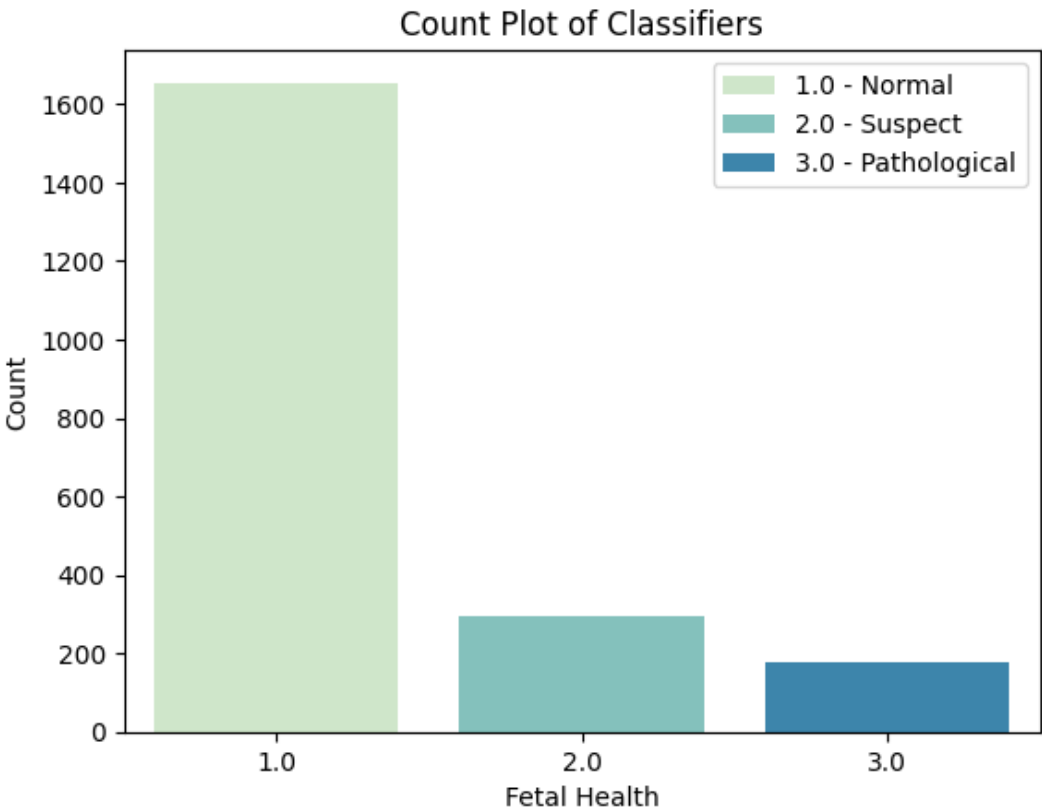
Step 2: Data Analysis

Classifier Distribution

In []:

```
labels = ['1.0 - Normal', '2.0 - Suspect', '3.0 - Pathological']
sns.countplot(data, x='fetal_health', palette='GnBu', label=labels)
plt.title('Count Plot of Classifiers')
plt.xlabel('Fetal Health')
plt.ylabel('Count')
```


```
plt.legend()  
plt.show()
```

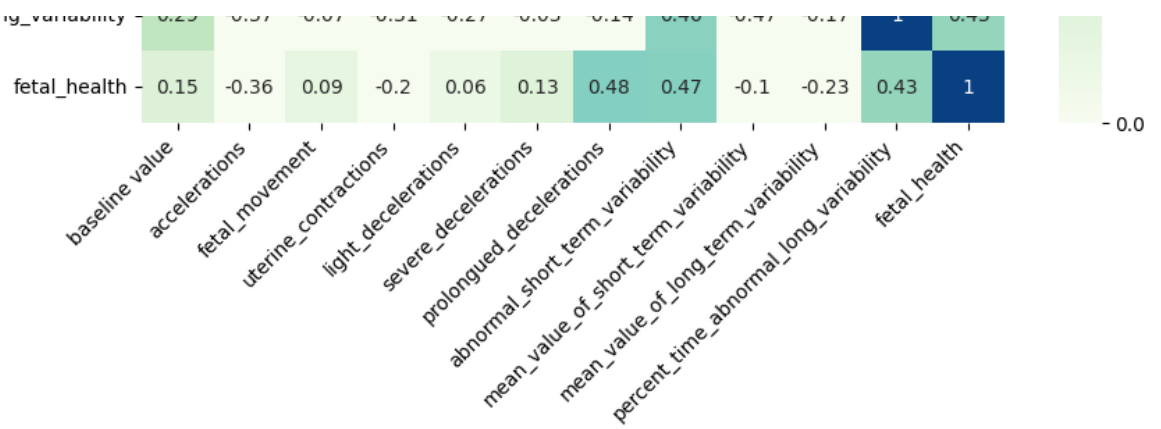


Correlation Matrix

In []:

```
corr = round(feat.corr(), 2)  
fig = plt.figure(figsize=(10,8))  
axes = fig.subplots()  
  
sns.heatmap(corr, vmin=0, vmax=1, annot=True, cmap='GnBu')  
  
plt.title('Correlation Matrix of Fetal Health Indicators')  
plt.xticks(ha='right', rotation=45)  
plt.show()
```

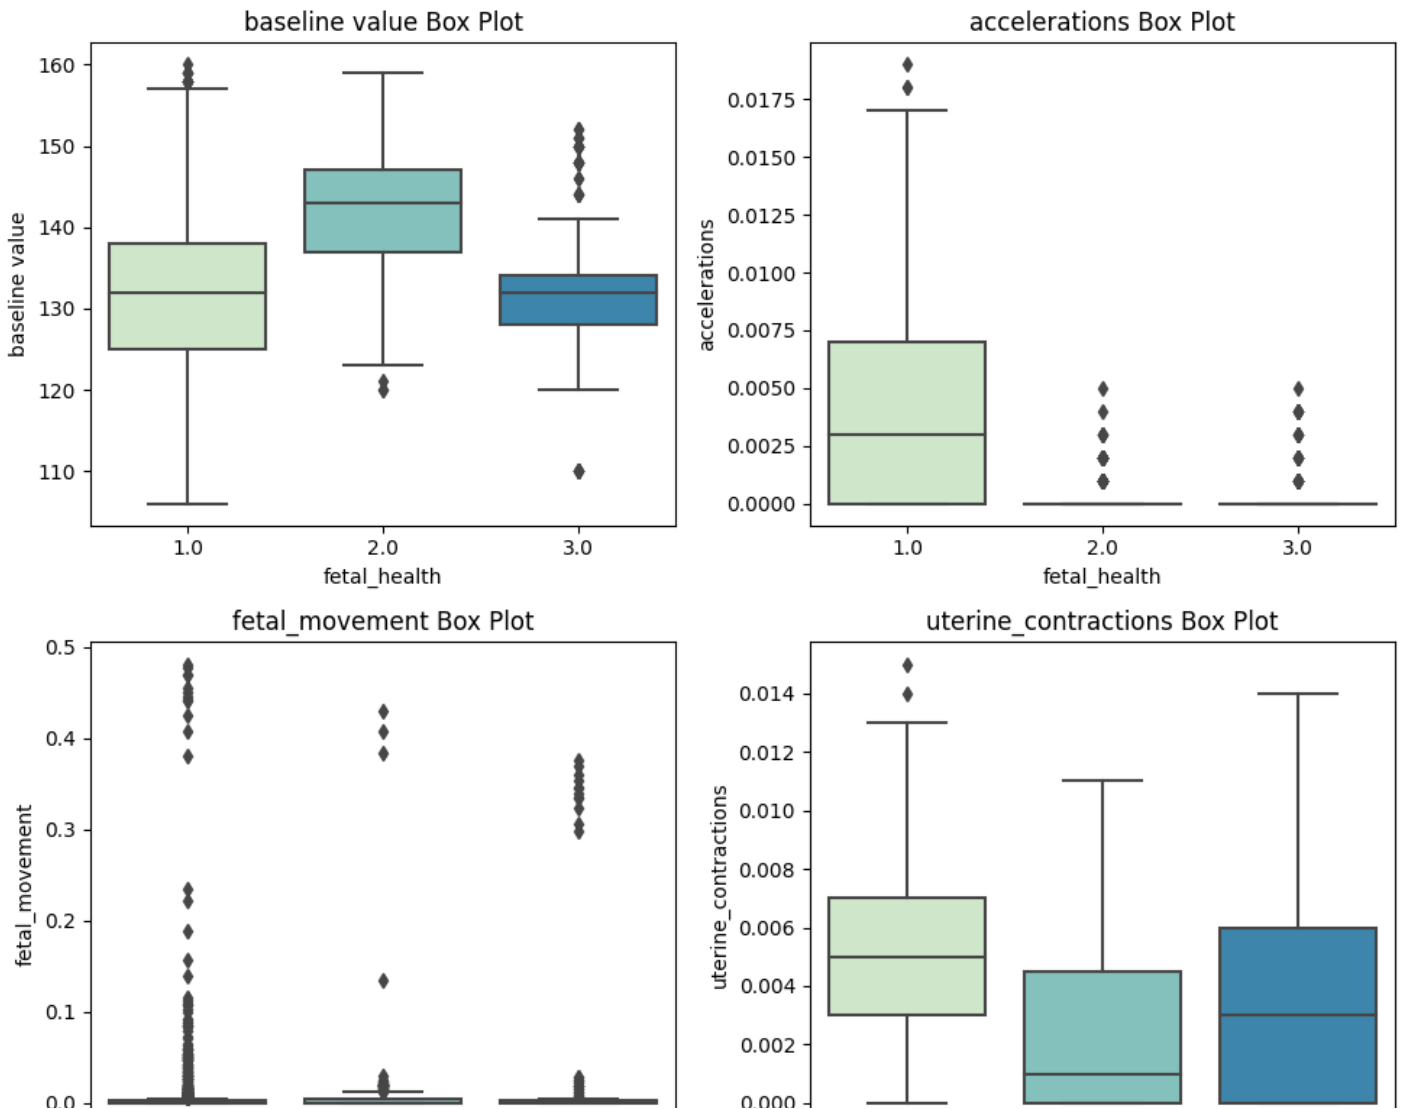



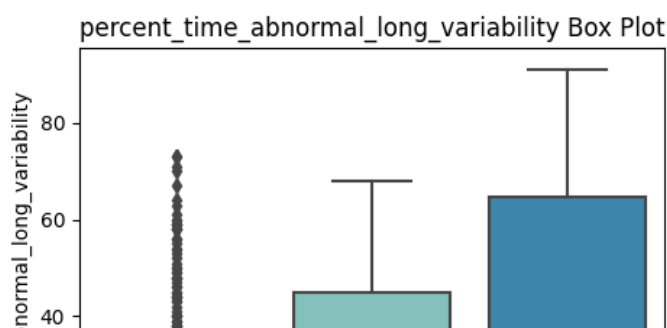
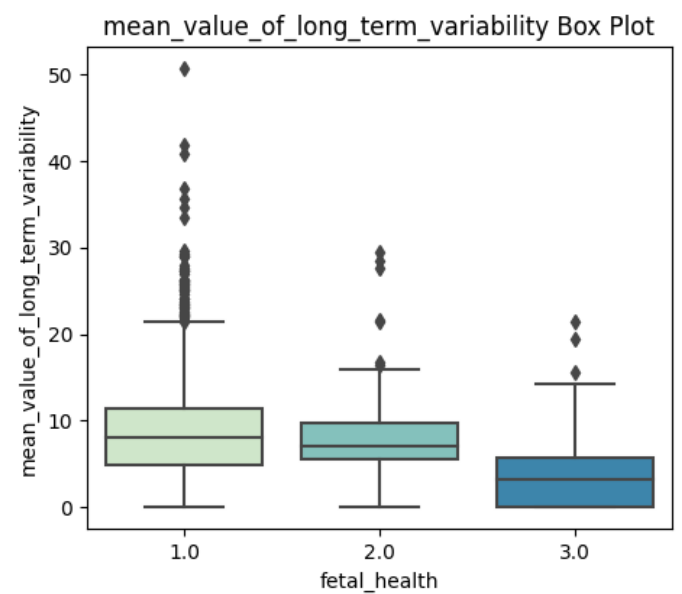

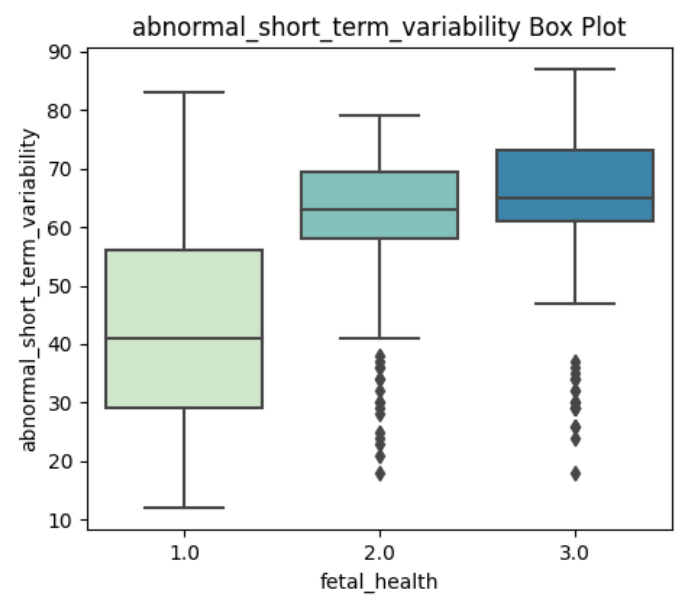
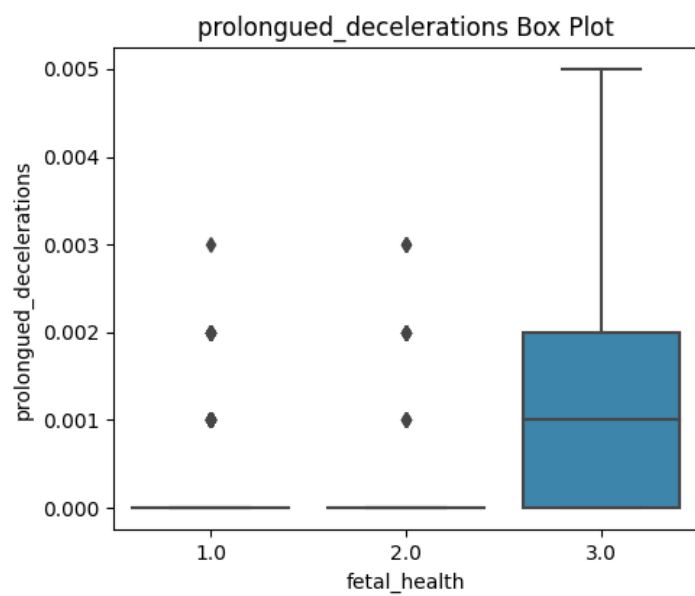
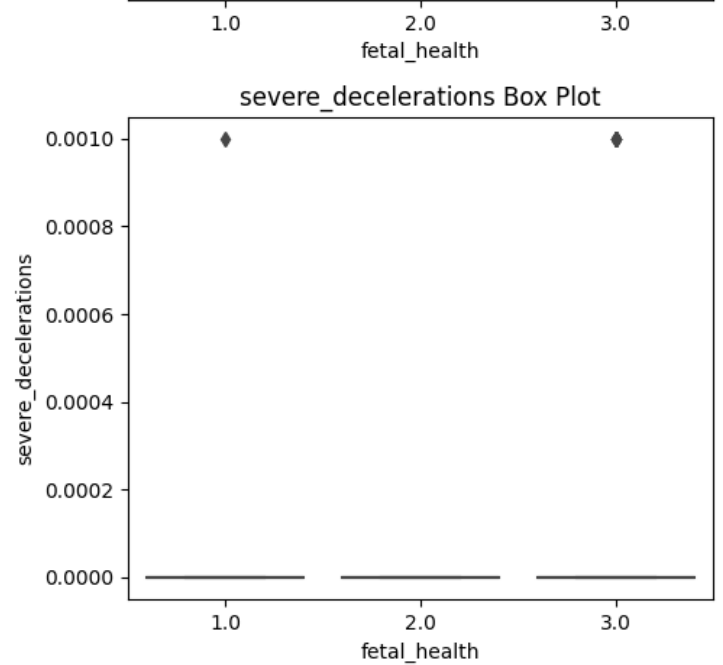
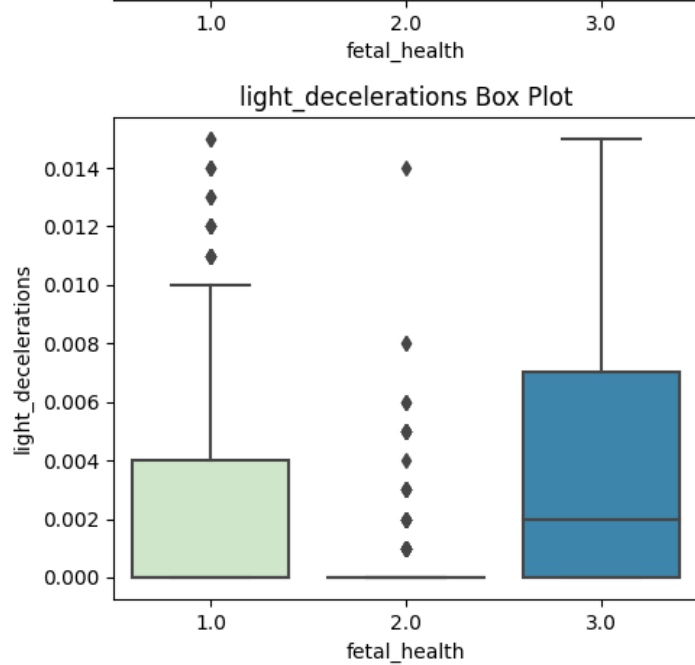




View Data Distribution - Box Plots

In []:

```
fig = plt.figure(figsize=(10,25)) # W x H
axes = fig.subplots(6,2)
temp = feat.columns.values
i=0
for row in range(6):
    for col in range(2):
        sns.boxplot(x=feat['fetal_health'],
                    y = feat[temp[i]],
                    data=feat,
                    ax=axes[row,col],
                    palette='GnBu')
        axes[row,col].set_title(f'{temp[i]} Box Plot')
        i+=1
fig.tight_layout()
plt.show()
```








View Data Distribution - Frequency Histograms

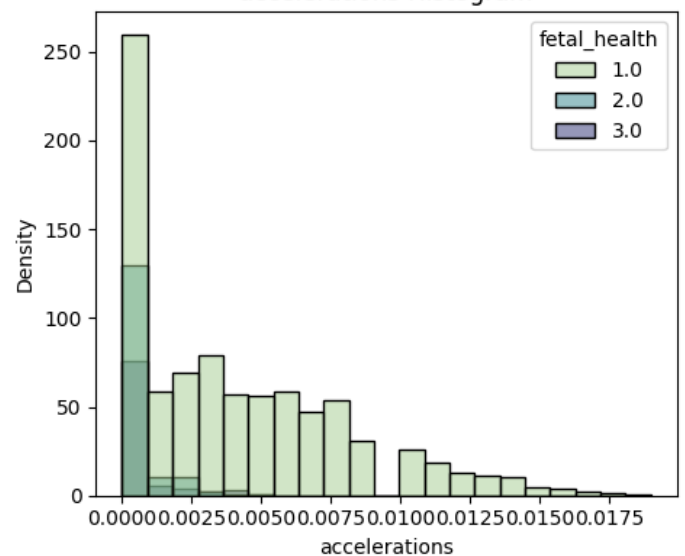
In []:

```
fig = plt.figure(figsize=(10,25)) # W x H
axes = fig.subplots(6,2)
temp = feat.columns.values
i=0
for row in range(6):
    for col in range(2):
        sns.histplot(data=feat,
                     x=feat[temp[i]],
                     hue=feat['fetal_health'],
                     stat='density',
                     palette='crest',
                     ax=axes[row,col],)
        axes[row,col].set_title(f'{temp[i]} Histogram')
    i+=1
fig.tight_layout()
plt.show()
```

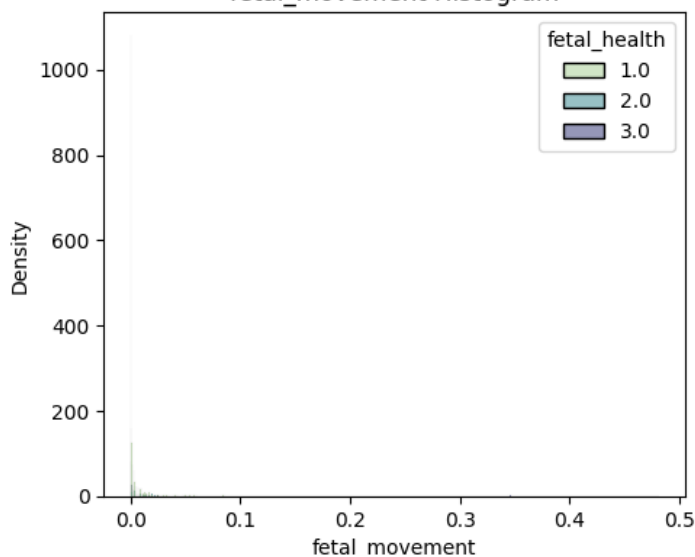
baseline value Histogram




accelerations Histogram



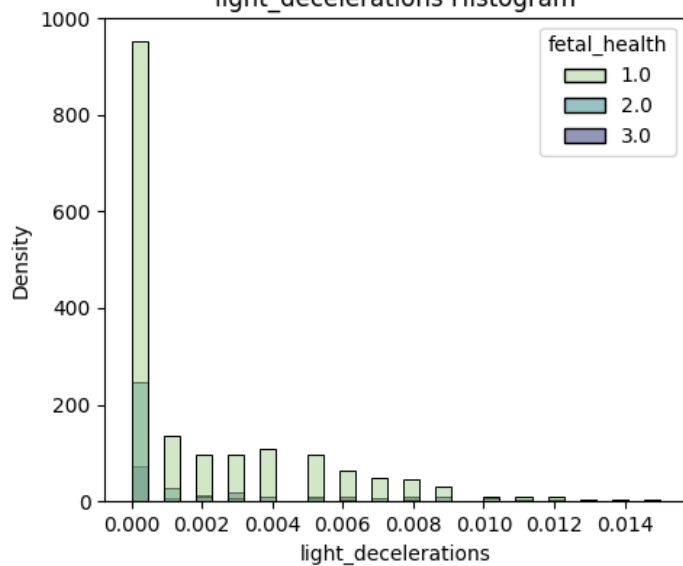
fetal_movement Histogram



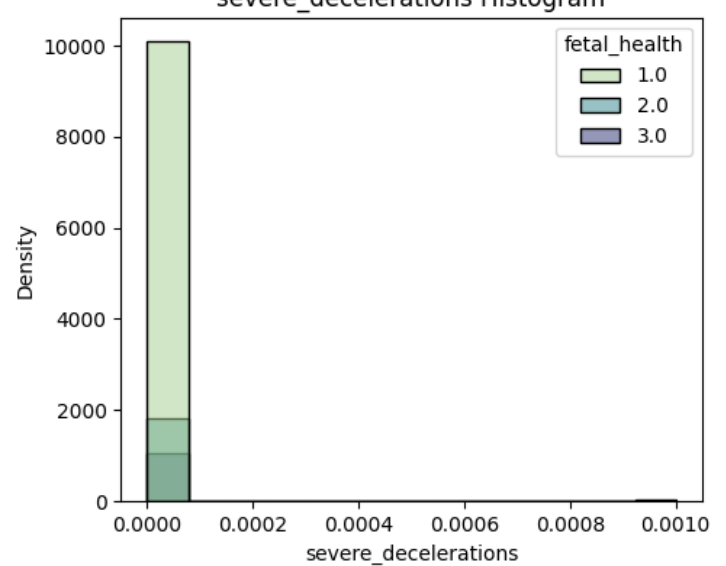
uterine_contractions Histogram




light_decelerations histogram




severe_decelerations histogram




prolonged_decelerations Histogram



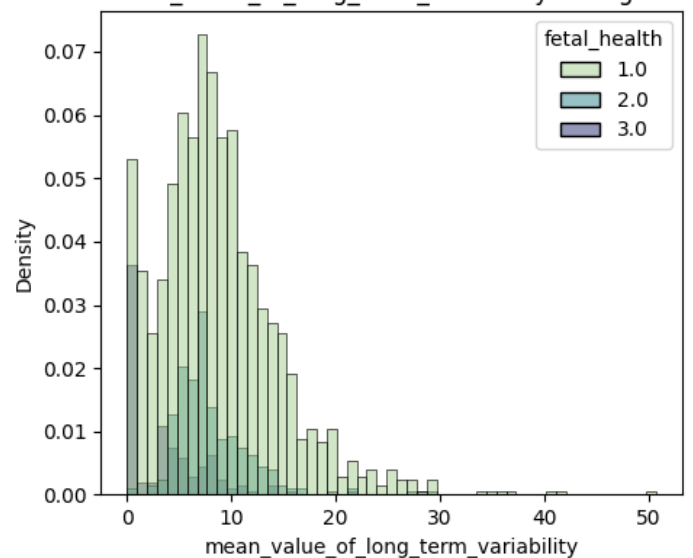
abnormal_short_term_variability Histogram




mean_value_of_short_term_variability Histogram



mean_value_of_long_term_variability Histogram





percent_time_abnormal_long_variability Histogram




fetal_health Histogram





In []:

```
from sklearn.preprocessing import MinMaxScaler
minmax = MinMaxScaler(feature_range=(0,1)).set_output(transform='pandas')
range_scaled = minmax.fit_transform(feats)
fig = plt.figure(figsize=(10,7))
sns.boxplot(range_scaled,
             palette='GnBu')
plt.title('Normalized Feature Box Plots')
plt.xticks(ha='right',rotation=45)
plt.show()
```



Step 3: Model Evaluation

In []:

```
from sklearn.model_selection import StratifiedShuffleSplit
from sklearn.model_selection import train_test_split
```

```

from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression

from sklearn.metrics import ConfusionMatrixDisplay
from sklearn.metrics import classification_report
from sklearn.inspection import DecisionBoundaryDisplay
from sklearn.preprocessing import LabelBinarizer
from sklearn.metrics import RocCurveDisplay
from sklearn.calibration import CalibratedClassifierCV
from sklearn.model_selection import GridSearchCV
from sklearn.svm import SVC

```

In []:

```

sc = StandardScaler().set_output(transform='pandas')
X_temp = feat.iloc[:, :-1]
X = sc.fit(X_temp).transform(X_temp).values
y = feat.iloc[:, -1].values
sss = StratifiedShuffleSplit(n_splits=2,
                             train_size=0.8,
                             test_size=0.2,
                             random_state=1234)

for i, (train_index, test_index) in enumerate(sss.split(X, y)):
    X_train = X[train_index]
    y_train = y[train_index]

    X_test = X[test_index]
    y_test = y[test_index]
print(X_train.shape)

label_binarizer = LabelBinarizer().fit(y_train)
y_onehot_test = label_binarizer.transform(y_test) # for ROC curves

(1700, 11)

```

In []:

```

def plot_split_results(data, param_dict, model_name):
    data = data.drop(columns=['mean_fit_time', 'std_fit_time', 'mean_score_time', 'std_score_time',
                             'mean_test_score', 'std_test_score', 'std_train_score',
                             'rank_test_score', 'split0_train_score', 'split1_train_score',
                             'split2_train_score', 'split3_train_score', 'split4_train_score',
                             'split5_train_score', 'split6_train_score', 'split7_train_score',
                             'split8_train_score', 'split9_train_score', 'mean_train_score'],)

    fig = plt.figure(figsize=(15,5))
    axes = fig.subplots(1, len(param_dict), sharey=True)

    for i, (param_name, param_range) in enumerate(param_dict.items()):
        fig.suptitle(f'10-Fold Cross-Validation Results for {model_name} Hyper-Parameters', fontsize=15)
        grouped = data.groupby(by=f'param_{param_name}').agg('mean').T
        grouped.index = np.arange(1, len(grouped) + 1)

        sns.lineplot(grouped, palette='ocean', ax=axes[i])
        axes[i].set(title=param_name,
                    xlabel='Fold',
                    ylabel='Score',
                    visible=True)
        axes[i].set_xticks(grouped.index)
        axes[i].legend()
    plt.show()

```

In []:

```
def plot_ROC_curve(mod, XX_train, yy_train, XX_test):
    fig = plt.figure(figsize=(5,10))
    axes=fig.subplots(2,1)
    ConfusionMatrixDisplay.from_predictions(y_true=y_test,
                                           y_pred=pred,
                                           cmap='GnBu',
                                           ax=axes[0])

    axes[0].set(title=f'{model_names[i]} Confusion Matrix')
    try:
        y_score = mod.fit(XX_train, yy_train).predict_proba(XX_test)
    except:
        m = CalibratedClassifierCV(mod)
        m.fit(XX_train, yy_train)
        y_score = m.predict_proba(XX_test)

    print(classification_report(y_test, pred))
    RocCurveDisplay.from_predictions(y_onehot_test.ravel(),
                                    y_score.ravel(),
                                    name="Micro-Average One-vs-Rest",
                                    color="navy",
                                    ax=axes[1])

    axes[1].set(title=f'{model_names[i]} ROC Curve', xlabel='False Positive Rate', ylabel=
'True Positive Rate')
    axes[1].plot([0, 1], [0, 1], label="chance level (AUC = 0.5)", color='lightgreen', lin
estyle='--')
    plt.legend()
    plt.show()
```

In []:

```
def plot_decision_boundary(mod):
    X_train_corr = X_temp[['abnormal_short_term_variability', 'percent_time_abnormal_long_v
ariability']]

    mod.fit(X_train_corr, y)

    disp = DecisionBoundaryDisplay.from_estimator(mod,
                                                X_train_corr,
                                                response_method="predict",
                                                alpha=0.5,
                                                xlabel='Abnormal Short Term Variability'
,
                                                ylabel='Percent Time Abnormal Long Varia
bility',
                                                cmap='GnBu')

    decision_function = mod.decision_function(X_train_corr)
    support_vector_indices = np.where(np.abs(decision_function) <= 1 + 1e-15)[0]
    support_vectors = X_train_corr.values[support_vector_indices]

    scat = plt.scatter(X_train_corr.iloc[:, 0],
                      X_train_corr.iloc[:, 1],
                      c=y,
                      edgecolors="k",
                      cmap='GnBu')
    plt.title('Decision Boundaries for Support Vector Machine')

    handles, labels = scat.legend_elements()
    labels = ['1.0 - Normal', '2.0 - Suspect', '3.0 - Pathological']
    plt.legend(handles=handles, labels=labels)

    plt.show()
```

In []:

```
model_names = ['Decision Tree',
               'Support Vector Machine',
               'Gradient Boost',
               'k-Nearest Neighbors',
```

```
'Logistic Regression']
```

```
dt = DecisionTreeClassifier(class_weight='balanced')
```

```
svm = SVC(class_weight='balanced')
```

```
gb = GradientBoostingClassifier()
```

```
knn = KNeighborsClassifier()
```

```
logreg = LogisticRegression(class_weight='balanced')
```

```
models = [dt, svm, gb, knn, logreg]
```

```
In [ ]:
```

```
dt_params = {'criterion':('gini', 'entropy', 'log_loss'),
             'max_features':('auto', 'sqrt', 'log2'),
             'splitter':('random', 'best')}
```

```
svm_params = {'C':(1.0, 1.25, 1.5, 1.75, 2.0),
              'kernel':('linear', 'rbf', 'poly', 'sigmoid'),
              'degree':(1,2,3,4,5)}
```

```
gb_params = {'criterion':('friedman_mse', 'squared_error'),
             'loss':('log_loss', 'deviance'),
             'n_estimators':(100, 150, 200, 250)}
```

```
knn_params = {'metric':('manhattan', 'euclidean', 'cosine'),
              'weights':('uniform', 'distance'),
              'n_neighbors':(3, 5, 7, 9)}
```

```
logreg_params = {'C':(1.0, 1.25, 1.5, 1.75, 2.0),
                 'penalty': ('l1', 'l2'),
                 'solver': ['liblinear']}
```

```
param_grids = [dt_params, svm_params, gb_params, knn_params, logreg_params]
```

```
In [ ]:
```


```
for i in range(len(models)):
    grid = GridSearchCV(estimator=models[i],
                       param_grid=param_grids[i],
                       scoring='f1_weighted',
                       return_train_score=True,
                       cv=10)
    grid.fit(X_train, y_train)
    results = pd.DataFrame(grid.cv_results_)#.loc[:, 'params': 'split9_test_score']
    plot_split_results(results, param_grids[i], model_names[i])
```


```
print(f'\n ----- {type(models[i]).__name__} ----- ')
print(f' The best estimator: {grid.best_estimator_}')
print(f' The best score: {round(grid.best_score_, 4)}')
print(f' The best parameters: {grid.best_params_}\n')
# display(results[['params', 'rank_test_score',
#                  'mean_test_score', 'std_test_score',
#                  'mean_train_score', 'std_train_score']])
```

```
model = grid.best_estimator_
pred = grid.predict(X_test)
```

```
plot_ROC_curve(model, X_train, y_train, X_test)
if model_names[i] == 'Support Vector Machine':
    plot_decision_boundary(model)
```

10-Fold Cross-Validation Results for Decision Tree Hyper-Parameters








```

----- DecisionTreeClassifier -----
The best estimator: DecisionTreeClassifier(class_weight='balanced', criterion='log_loss',
                                          max_features='log2')
The best score: 0.9153
The best parameters: {'criterion': 'log_loss', 'max_features': 'log2', 'splitter': 'best'}


```

	precision	recall	f1-score	support
1.0	0.93	0.94	0.94	332
2.0	0.70	0.68	0.69	59
3.0	0.76	0.74	0.75	35
accuracy			0.89	426
macro avg	0.80	0.79	0.79	426
weighted avg	0.89	0.89	0.89	426





10-Fold Cross-Validation Results for Support Vector Machine Hyper-Parameters



----- SVC -----


The best estimator: SVC(class_weight='balanced', kernel='poly')

The best score: 0.9091

The best parameters: {'C': 1.0, 'degree': 3, 'kernel': 'poly'}


	precision	recall	f1-score	support
1.0	0.96	0.93	0.94	332
2.0	0.71	0.81	0.76	59
3.0	0.78	0.83	0.81	35
accuracy			0.90	426
macro avg	0.82	0.86	0.83	426
weighted avg	0.91	0.90	0.91	426

Support Vector Machine Confusion Matrix




Support Vector Machine ROC Curve






Decision Boundaries for Linear Support Vector Machine






10-Fold Cross-Validation Results for Gradient Boost Hyper-Parameters




----- GradientBoostingClassifier -----

The best estimator: GradientBoostingClassifier(loss='deviance')
The best score: 0.9376
The best parameters: {'criterion': 'friedman_mse', 'loss': 'deviance', 'n_estimators': 100}



	precision	recall	f1-score	support
1.0	0.95	0.98	0.97	332
2.0	0.85	0.75	0.79	59
3.0	0.94	0.89	0.91	35
accuracy			0.94	426
macro avg	0.91	0.87	0.89	426
weighted avg	0.94	0.94	0.94	426







----- KNeighborsClassifier -----
 The best estimator: KNeighborsClassifier(metric='manhattan', n_neighbors=7, weights='distance')
 The best score: 0.9293
 The best parameters: {'metric': 'manhattan', 'n_neighbors': 7, 'weights': 'distance'}

	precision	recall	f1-score	support
1.0	0.94	0.96	0.95	332
2.0	0.71	0.66	0.68	59
3.0	0.88	0.83	0.85	35
accuracy			0.91	426
macro avg	0.84	0.82	0.83	426
weighted avg	0.91	0.91	0.91	426





10-Fold Cross-Validation Results for Logistic Regression Hyper-Parameters




----- LogisticRegression -----

The best estimator: LogisticRegression(class_weight='balanced', solver='liblinear')


The best score: 0.8765


The best parameters: {'C': 1.0, 'penalty': 'l2', 'solver': 'liblinear'}

	precision	recall	f1-score	support
1.0	0.96	0.91	0.93	332
2.0	0.69	0.75	0.72	59
3.0	0.62	0.80	0.70	35
accuracy			0.88	426
macro avg	0.76	0.82	0.78	426
weighted avg	0.89	0.88	0.88	426



Logistic Regression ROC Curve





In []: